

Blotch3D User Manual

Create real-time 3D graphics with just a few lines of C# code.

Quick start

Introduction

Project structure

Development

Making 3D models

Dynamically changing a sprite's orientation and position

Matrix internals

A Short Glossary of 3D Graphics Terms

Troubleshooting

Rights

Quick start

1. Get the installer for the latest release of MonoGame from <http://www.monogame.net/downloads/> and run it. (Do NOT get the current development version nor the NuGet package.)
2. Get the Blotch3D repository zip from <https://github.com/Blotch3D/Blotch3D> and unzip it.
3. Open the Visual Studio solution file (Blotch3D.sln).
4. Build and run the example projects. (For other platforms, you'll need the appropriate Visual Studio add-on and you will need to create a separate project for that platform.)
5. Use IntelliSense to see the reference documentation, or view the Doxygen output of it in "Blotch3D_Reference_Manual.pdf".

Introduction

Blotch3D is a C# library that vastly simplifies many of the tasks in developing 3D applications and games.

Examples are provided that show how with just a few lines of code you can...

- Load standard 3D model file types as “sprites”, and display and move thousands of them in 3D at high frame rates.
- Set a sprite’s material, texture, and lighting response.
- Load textures from standard image files.
- Show 2D and in-world (as a texture) text in any font, size, color, etc. at any 2D or 3D position, and make text follow a sprite in 2D or 3D.
- Attach sprites to other sprites to create ‘sprite trees’ as large as you want. Child sprite orientation and position is relative to its parent sprite’s orientation and position, and can be changed dynamically (i.e. the sprite trees are dynamic scene graphs.)
- Override all steps in the drawing of each sprite.
- You can give the user easy control over all aspects of the camera (zoom, pan, truck, dolly, rotate, etc.).
- Easily control all aspects of the camera programmatically.
- Create billboard sprites.
- Connect sprites to the camera to implement HUD models and text.
- Connect the camera to a sprite to implement ‘cockpit view’, etc.
- Implement GUI controls (as dynamic 2D text or image rectangles) in the 3D window.
- Implement a skybox.
- Get a list of sprites touching a ray, to implement weapons fire, etc.
- Get a list of sprites under the mouse position, to implement mouse selection, tooltips, pop-up menus, etc.
- Implement levels-of-detail.
- Implement mipmaps.
- Implement translucent sprites and textures with an alpha channel.
- Create sprite models programmatically (custom vertices).
- Use with WPF and WinForms, on Microsoft Windows.
- Access and override many window features and functions using the provided WinForms Form object of the window (Microsoft Windows only).
- Detect collisions between sprites.
- Implement fog
- Define ambient lighting, and up to three point-light sources. (More lights can be defined if a custom shader is used.)

- Build for many platforms (currently supports all Microsoft Windows platforms, iOS, Android, MacOS, Linux, PS4, PSVita, Xbox One, and Switch).

Blotch3D sits on top of MonoGame. MonoGame is a widely used 3D library for C#. It is free, fast, cross platform, actively developed by a large community, and it's used in many professional games. There is a plethora of MonoGame documentation, tutorials, examples, and discussions on line. All MonoGame features remain available. For example, custom shaders can be written to override the default shader.

All reference documentation of Blotch3D (classes, methods, fields, properties, etc.) is available through Visual Studio IntelliSense, and it is also available in Doxygen output file “Blotch3D_Reference_Manual.pdf”. (Note: To support Doxygen, links in the IntelliSense comments are preceded with ‘#’. Also, this document is prepended to that Doxygen output, but for some reason the reference links within this document don't work when it is embedded in that reference manual.)

See MonoGame.net for the official MonoGame documentation. When searching on-line for other MonoGame documentation and discussions, be sure to note the MonoGame version being discussed. Documentation of earlier versions may not be compatible with the latest.

MonoGame fully implements Microsoft's (no longer supported) XNA 4 engine, but for multiple platforms. It also implements features beyond XNA 4. Therefore XNA 4 documentation you come across may not show you the best way to do something, and documentation of earlier versions of XNA (versions 2 and 3) will often not be correct. For conversion of XNA 3 to XNA 4 see <http://www.nelsonhurst.com/xna-3-1-to-xna-4-0-cheatsheet/>.

Note that to support all the platforms, certain limitations were necessary. Currently you can only have one 3D window. Also, there is no official cross-platform way to specify an existing window to use as the 3D window—MonoGame must create it. See below for details and work-arounds.

Project structure

The provided Visual Studio solution file contains both the Blotch3D library project with source, and the example projects.

“BlotchExample01_Basic” is a bare-bones Blotch3D application, where Example.cs contains the example code. Other example projects also contain an Example.cs, which is similar to the one from the basic example but with a few additions to it to demonstrate a certain feature. In fact, you can do a diff between the “BlotchExample01_Basic” source file and another example's source file to see what extra code must be added to implement the features it demonstrates [TBD: the “full” example needs to be split to several simpler examples].

All the provided projects are configured to build for the Microsoft Windows x64 platform. See below for other platforms.

To create a new project, you can just copy the basic example and rename the project, or you can create the project from scratch like this:

1. If you haven't already done it, install MonoGame as described in the Quick start section.
2. If you are building for a platform other than Microsoft Windows, install the Visual Studio add-ons, etc. for that platform. (For example, for Android you'd need Xamarin for Android.)
3. Create a new project for a platform that is compatible with MonoGame.
4. Add a reference to MonoGame. (For .NET Framework, you would add something like `\Program Files (x86)\MonoGame\v3.0\Assemblies\Windows\MonoGame.Framework.dll`)
5. If the Blotch3D project is not in the solution, add a reference to the Blotch3D assembly (like `Blotch3D.dll` on Microsoft Windows).
6. Follow the procedure in the 'Making 3D models' section to add a content folder and the pipeline manager so that you have a way to add content.
7. You'll probably want to set the output type to 'Console Application' for now, so you can see any debug messages. You can change this to 'Windows Application' later, if you like.
8. To create a 3D window, follow the guidelines in the Development section.

The above process works for Microsoft Windows. Generally this should be close to the procedure you need for other platforms.

If you are copying the Blotch3D assembly (like `Blotch3D.dll` on Microsoft Windows) to a project or packages folder so you don't have to include the source code of the library in your solution, be sure to also copy `Blotch3D.xml` so you still get the IntelliSense. You shouldn't have to copy any other binary file from the Blotch3D output folder if you've installed MonoGame on the destination machine. Otherwise you should copy the entire project output folder. For example, you'd probably want to copy everything in the output folder when you are distributing your app.

Development

See the examples.

To make a 3D window, you must derive a class from `BlWindow3D` and override the `Setup`, `FrameProc`, and `FrameDraw` methods.

When it comes time to open the 3D window, you instantiate that class and call its "Run" method *from the same thread that instantiated it*. The Run method

will call the Setup, FrameProc, and FrameDraw methods when appropriate (explained below), and not return until the window closes.

The thread that instantiates the BIWindow3D-derived class, calls the Run method, etc., we will call the “3D thread”.

All code that accesses 3D resources must be done in the 3D thread, including code that creates and uses all Blotch3D and MonoGame objects. Note that this rule also applies to any code structure that may internally use other threads, as well. Do not use Parallel, async, etc. code structures that access 3D resources.

This pattern and these rules are also used by MonoGame. In fact, the BIWindow3D class inherits from MonoGame’s “Game” class. But instead of overriding certain “Game” class methods, you override BIWindow3D’s Setup, FrameProc, and FrameDraw methods. Other “Game” class methods and events can still be overridden, if needed.

All this is necessary because certain 3D subsystems (OpenGL, DirectX, etc.) generally require that 3D resources be accessed by a single thread. (There are some platform-specific exceptions, but MonoGame does not use them.)

The Setup, FrameProc, and FrameDraw override methods are called by the 3D thread as follows:

The Setup method is called by the 3D thread once at the beginning of instantiation of the BIWindow3D-derived object. You might put time-consuming initialization of persistent things in there like loading of persistent content (sprite models, fonts, etc.), creation of persistent BLSprites, etc.

The FrameProc method is called by the 3D thread once per frame. For single-threaded applications this is typically where the bulk of application code resides, except the actual drawing code. For multi-threaded applications, this is where all application code resides that does anything with 3D resources.

The FrameDraw method is called by the 3D thread every frame, but only if there is enough CPU for the thread. Otherwise it calls it less frequently. This is where you put drawing code (BLSprite.Draw, BLSprite.DrawText, etc. You can also put app code here as long as it’s aware that calls to it may not be periodic.

A single-threaded application would have all its code in those three overridden methods.

If you are developing a multithreaded app, then when other threads need to create, change, or destroy 3D resources or otherwise do something in a thread-safe way with the 3D thread, they can pass a delegate to EnqueueCommand or EnqueueCommandBlocking. Those methods make sure the code is done by the 3D thread sequentially at the end of the next FrameProc call.

MonoGame does not support multiple 3D windows because that isn’t conducive on certain platforms. On Microsoft Windows (and possibly certain other

platforms) you *can* create them, but they don't work correctly and in certain situations will crash. If you want to be able to “close” and “re-open” a window, you can just hide and show the same window. (On Microsoft Windows, you can use the `BlWindow3D.Form` object for that.)

To make the MonoGame window be a child window of an existing GUI, you need to explicitly size, position, and convey Z order so that it is overlaid at the right screen location. The easiest way to do that would be to overlay the 3D window on an existing child window by getting the current attributes of that child window, whenever they change. On Microsoft Windows, the window's Form object (`BlWindow3D.Form`) may be of help in this. There may also be a way to specify that an existing window be used as the 3D window, but it probably isn't portable and may not work in later MonoGame releases.

Most Blotch3D objects must be Disposed when you are done with them and you are not otherwise terminating the program.

See the examples and use IntelliSense for more information.

Making 3D models

There are several primitive models available with Blotch3D. The easiest way to add them to your project is to...

1. Copy the Content folder from the Blotch3D project folder to your project folder
2. Add the “Content.mgcb” file in that folder to your project
3. Right-click it and select “Properties”
4. Set the “Build Action” to “MonoGameContentReference”

You can get the names of the content files by starting the MonoGame pipeline manager (double-click Content/Content.mgcb). You can also add more content via the pipeline manager (see <http://rbwhitaker.wikidot.com/monogame-managing-content>). See the examples for details on how to load and display models, fonts, etc.

If no existing model meets your needs, you can either programmatically create a model by specifying the vertices and normals (see the example that uses custom Vertices), or create a model with, for example, the Blender 3D modeler and then add it to the project with the pipeline manager. The pipeline manager can import several model file types. You can also instruct Blender to include texture (UV) mapping by using one of the countless tutorials online, like <https://www.youtube.com/watch?v=2xTzJIaKQFY> or https://en.wikibooks.org/wiki/Blender_3D:_Noob_to_Pro/UV_Map_Basics. Also, you may be able to import certain existing models from the web.

Dynamically changing a sprite's orientation and position

Each sprite has a “Matrix” member that defines its orientation and position relative to its parent sprite, or to an unmodified coordinate system if there is no parent. There are many static and instance methods of the Matrix class that let you easily set and change the scaling, translation, rotation, etc. of a matrix.

When you change anything about a sprite's matrix, you also change the orientation and position of its child sprites, if any. That is, subsprites reside in the parent sprite's coordinate system. For example, if a child sprite's matrix scales it by 3, and its parent sprite's matrix scales by 4, then the child sprite will be scaled by 12. Likewise, rotation, shear, and translation are inherited, as well.

There are also static and instance Matrix methods and operator overloads to combine (multiply) matrices to form a single matrix which combines the effects of multiple matrices. For example, a rotate matrix and a scale matrix can be multiplied to form a single rotate-scale matrix. But mind the multiplication order because matrix multiplication is not commutative. See below for details, but novices can simply try the operation one way (like A times B) and, if it doesn't work the way you wanted, do it the other way (B times A).

For a good introduction (without the math), see <http://rbwhitaker.wikidot.com/monogame-basic-matrices>.

The rest of this section should be studied only when you need a deeper knowledge.

Matrix internals

Here we'll introduce the internals of 2D matrices. 3D matrices simply have one more dimension.

Let's imagine a model that has one vertex at (4,1) and another vertex at (3,3). (This is a very simple model comprised of only two vertices!)

You can move the model by moving each of those vertices by the same amount, and without regard to where each is relative to the origin. To do that, just add an offset vector to each vertex. For example, we could add the vector (2,1) to each of those original vertices, which would result in final model vertices of (6,2) and (5,4). In that case we have *translated* (moved) the model.

Matrices certainly support translation. But first let's talk about moving a vertex *relative to its current position from the origin*, because that's what gives matrices the power to shear, rotate, and scale a model about the origin. This is because those operations affect each vertex differently depending on its relationship to the origin.

If we want to scale (stretch) the X relative to the origin, we can multiply the X of each vertex by 2.

For example,

$X' = 2X$ (where X is the initial value, and X' is the final value)

... which, when applied to each vertex, would change the above vertices from (4,1) and (3,3) to (8,1) and (6,3).

We might want to define how to change each X according to the original X value of each vertex *and also according to the original Y value*, like this:

$$X' = aX + bY$$

For example, if $a=0$ and $b=1$, then this would set the new X of each vertex to its original Y value.

Finally, we might also want to define how to create a new Y for each vertex according to its original X and original Y . So, the equations for both the new X and new Y are:

$$X' = aX + bY$$

$$Y' = cX + dY$$

(Remember, the idea is to apply this to every vertex.)

By convention we might write the four matrix elements (a , b , c , and d) in a 2x2 matrix, like this:

a b

c d

This should all be very easy to understand.

But why are we even talking about it? Because now we can define the elements of a matrix that, if applied to each vertex of a model, define any type of *transform* in the position and orientation of that model.

For example, if we apply the following matrix to each of the model's vertices:

1 0

0 1

... then the vertices are unchanged, because...

$$X' = 1X + 0Y$$

$$Y' = 0X + 1Y$$

... sets X' to X and Y' to Y .

This matrix is called the *identity* matrix because the output (X', Y') is the same as the input (X, Y).

We can create matrices that scale, shear, and even rotate points. To make a model three times as large (relative to the origin), use the matrix:

3 0

0 3

To scale only X by 3 (stretch a model in the X direction about the origin), then use the matrix:

3 0

0 1

The following matrix flips (mirrors) the model vertically about the origin:

1 0

0 -1

Below is a matrix to rotate a model counterclockwise by 90 degrees about the origin:

0 -1

1 0

Here is a matrix that rotates a model counterclockwise by 45 degrees about the origin:

0.707 -0.707

0.707 0.707

Note that '0.707' is the sine of 45 degrees.

A matrix can be created to rotate any amount about any axis.

(The Matrix class provides functions that make it easy to create a rotation matrix from a rotation axis and angle, or pitch and yaw and roll, or something called a quaternion, since otherwise we'd have to call sine and cosine functions, ourselves, to create the matrix elements.)

Since we often also want to translate (move) points *without* regard to their current distances from the origin as we did at the beginning of this section, we add more numbers to the matrix just for that purpose. And since many mathematical operations on matrices work only if the matrix has the same number of rows as columns, we add more elements simply to make the rows and columns the same size. And since Blotch3D/MonoGame works in 3-space, we add even more numbers to handle the Z dimension. So, the final matrix size in 3D graphics is 4x4.

Specifically:

$$X' = aX + bY + cZ + d$$

$$Y' = eX + fY + gZ + h$$

$$Z' = iX + jY + kZ + l$$

$$W = mX + nY + oZ + p$$

(Consider the W as unused, for now.)

Which can be notated as...

a b c d

e f g h

I j k l

m n o p

Notice that the d, h, and l are the translation vector.

The Matrix class in MonoGame uses the following field names:

M11 M12 M13 M14

M21 M22 M23 M24

M31 M32 M33 M34

M41 M42 M43 M44

Besides the ability to multiply entire matrices (as mentioned at the beginning of this section), you can also divide (i.e. multiply by a matrix inverse) matrices to, for example, solve for a matrix that was used in a previous matrix multiply, or otherwise isolate one operation from another. Welcome to linear algebra! We won't get in to how matrix multiplication and division specifically process the individual elements of the matrices because the Matrix class already provides those static or instance functions.

As was previously mentioned, each sprite has a matrix describing how that sprite and its children are transformed from the parent sprite's coordinate system. Specifically, Blotch3D does a matrix-multiply of the parent's matrix by the child's matrix to create the final matrix used to draw that child, and it is also used as the parent matrix for the subsprites of that child.

Because of confusion in coordinate system handedness (chirality), multiplication/division order, row vs. column notation (mathematicians use the opposite notation of that used by 3D graphics people), and the order of element storage in memory; on occasion it may be easier to try things one way and, if it doesn't work as expected, try it another way. But for details see <http://seanmiddleditch.com/matrices-handedness-pre-and-post-multiplication-row-vs-column-major-and-notations>.

A Short Glossary of 3D Graphics Terms

Vertex

A point in space. Typically, a point at which the line segments of a polygon meet.

That is, a corner of a polygon. A corner of a model. Most visible models are described as a set of vertices. Each vertex can have a color, texture coordinate, and normal.

Polygon

A visible surface described by a set of vertices that define its corners. A triangle is a polygon with three vertices, a quad is a polygon with four. One side of a polygon is a "face".

Ambient lighting

A 3D scene has one ambient light setting. The intensity of ambient lighting on the surface of a polygon is unrelated to the orientation of the polygon or the camera.

Diffuse lighting

Directional or point source lighting. You can have multiple directional or point light sources. Its intensity depends on the orientation of the polygon relative to the light.

Texture

A 2D image applied to the surface of a model. For this to work, each vertex of the model must have a texture coordinate associated with it, which is an X,Y coordinate of the 2D bitmap image that should be aligned with that vertex. Pixels across the surface of a polygon are interpolated from the texture coordinates specified for each vertex.

Normal

In mathematics, the word "normal" means a vector that is perpendicular to a surface. In 3D graphics, "normal" means a vector that indicates from what direction light will cause a surface to be brightest. Normally they would mean the same thing. However, by defining a normal at some angle other than perpendicular, you can somewhat cause the illusion that a surface lies at a different angle. Each vertex of a polygon has a normal vector associated with it and the brightness across the surface of a polygon is interpolated from the normals of its vertices. So, a single flat polygon can have a gradient of brightness across it giving the illusion of curvature. In this way a model composed of fewer polygons can still be made to look quite smooth.

X-axis

The axis that extends right from the origin.

Y-axis

The axis that extends forward from the origin.

Z-axis

The axis that extends up from the origin.

Translation

Movement. The placing of something at a different location from its original location.

Rotation

The circular movement of each vertex of a model about the same axis.

Scale

A change in the width, height, and/or depth of a model.

Shear (skew)

A pulling of one side of a model in one direction, and the opposite side in the opposite direction, without rotation, such that the model is distorted rather than rotated. A parallelogram is a rectangle that has experienced shear. If you apply another shear along an orthogonal axis of the first shear, you rotate the model.

Yaw

Rotation about the Y-axis

Pitch

Rotation about the X-axis, after any Yaw has been applied.

Roll

Rotation about the Z-axis, after any Pitch has been applied.

Euler angles

The yaw, pitch, and roll of a model, applied in that order.

Matrix

An array of 16 numbers that describes the position and orientation of a sprite. Specifically, a matrix describes a difference, or transform, in the orientation (coordinate system) of one model from another. See Dynamically changing a sprite's orientation and position.

Origin

The center of a coordinate system. The point in the coordinate system that is, by definition, at (0,0).

Frame

In this document, 'Frame' means a complete still scene. It is analogous to a movie frame. A moving 3D scene is created by drawing successive frames—typically at about 15 to 60 times per second.

Depth buffer

3D systems must keep track of the depth of the polygon surface (if any) at each 2D pixel so that they know to draw the nearer pixel over the farther pixel in the 2D display. The depth buffer is an array with one element per 2D screen pixel, where each element is (typically) a 32-bit floating point value indicating the depth of the nearest polygon surface at that pixel. See 'near clip' and 'far clip'.

Near clipping plane (near clip)

The distance from the camera at which a depth buffer element is equal to zero. Nearer surfaces are not drawn.

Far clipping plane (far clip)

The distance from the camera at which a depth buffer element is equal to the maximum possible floating-point value. Farther surfaces are not drawn.

Troubleshooting

Q: When I set a billboard attribute of a flat sprite (like a plane), I can no longer see it.

A: Perhaps the billboard orientation is such that you are looking at the plane from the side or back. Try setting a rotation in the sprite's matrix (and make sure it doesn't just rotate it on the axis intersecting your eye point).

Q: When I'm inside a sprite, I can't see it.

A: By default, Blotch3D draws only the outside of a sprite. Try doing a "Graphics.GraphicsDevice.RasterizerState = RasterizerState.CullClockwise" (or set it to CullNone to see both the inside and outside) in the BLSprite.PreDraw delegate, and set it back to CullCounterClockwise in the BLSprite.DrawCleanup delegate.

Q: I set a sprite's matrix so one of the dimensions has a scale of zero, but then the sprite becomes black.

A: A sprite's matrix also affects its normals. By setting a dimension's scale to zero, you may have caused some of the normals to be zero'd out as well.

Q: When I am zoomed-in a large amount, sprite and camera movement jumps.

A: You are experiencing floating point precision errors in the positioning algorithms. About all you can do is "fake" being that zoomed in by, instead, moving the camera forward temporarily. Or simply don't allow zoom to go to that extreme.

Q: Sometimes I see the polygons and parts polygons of sprites appear and disappear randomly as the camera or sprite moves.

A: The floating-point precision limitation of the depth buffer can cause this. Try increasing your near clip and/or decreasing your far clip so the depth buffer doesn't have to cover so much dynamic range.

Q: I have a sprite that I want to always be visible, but I think its invisible because its outside the depth buffer.

A: Try doing a "Graphics.GraphicsDevice.DepthStencilState = Graphics.DepthStencilState.Disabled" in the BLSprite.PreDraw delegate, and set it back to DepthStencilState.Enabled in the BLSprite.DrawCleanup delegate.

Rights

Blotch3D (formerly GWin3D) is Copyright © 1999-2018 by Kelly Loum

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.